

Querying a MySQL database

This lab was originally created by Profs Ben Baumer and Jordan Crouser. I've edited it slightly.

SQL is a longstanding database querying language. It is a loosely-implemented standard. We will be using MySQL.

To facilitate our connection to the MySQL database server, we will need to install the RMySQL package.

```
# do NOT install RMySQL if you are on the RStudio Server
install.packages("RMySQL")
```

Goal: by the end of this lab, you will be able to write basic SELECT queries in SQL and retrieve the results into R.

Connecting to MySQL

The data we will be using is stored on a server in Bass Hall. It's called `scidb.smith.edu`. We can connect through the `dbConnect()` function provided by the DBI package (which is loaded automatically when you load RMySQL). You will also need the RMySQL package installed.

```
library(tidyverse)
library(RMySQL)
db <- dbConnect(
  MySQL(),
  host = "scidb.smith.edu",
  user = "sds192",
  password = "DSismfc@S",
  dbname = "imdb"
)
knitr::opts_chunk$set(connection = 'db', max.print = 5)
```

This chunk of code will allow you to connect to `scidb`. Note that this creates a database connection object named `db`, which has the class `MySQLConnection`.

```
class(db)
```

```
[1] "MySQLConnection"  
attr(,"package")  
[1] "RMySQL"
```

Also, we set the `connection` parameter for all future chunks in this R Markdown file. Note also that the `max.print` argument sets the maximum number of results printed by each query.

Each of the following chunks makes use of the SQL engine functionality in `knitr`. You may want to [read about this](#). Each of the following chunks is an `sql` chunk – as opposed to an `r` chunk!

To retrieve the results from a query in R, use the `dbGetQuery()` function from the DBI package (which is automatically loaded when you load `RMySQL`). Its first argument is a database connection object, and the second argument is an SQL query as a character vector.

Retrieving data

We want to be able to see which type of databases exist on the server. We can do this with `SHOW DATABASES`

```
SHOW DATABASES;
```

Table 1: Displaying records 1 - 5

Database
airlines
citibike
fec
imdb
information_schema

You don't actually need the `;` at the end of the query above for MySQL, but for other SQL dialects you do so it doesn't hurt.

We want to use the `imdb` database and we have to tell our `db` connection that.

```
USE imdb;
```

It's fine that there are 0 rows, this is just how we tell R which db we want to use.

Let's look at the tables available in imdb.

```
SHOW TABLES;
```

Table 2: Displaying records 1 - 5

Tables_in_imdb
aka_name
aka_title
cast_info
char_name
comp_cast_type

See the kind_type table? That one shows what type of movie it is.

This query returns the list of kinds of “movies” stored in the IMDB. We are selecting everything with * from the kind_type table.

```
SELECT * FROM kind_type;
```

Table 3: 7 records

id	kind
1	movie
2	tv series
3	tv movie
4	video movie
5	tv mini series
6	video game
7	episode

Of course, you will often want to store the result of your query as a data frame. This can be achieved by setting the `output.var` argument in the chunk. Here we retrieve the list of different types of information stored in the database, save it as a data frame in R, and show the first few rows.

```
## output.var="info_types" in this chunk
```

```
SELECT * FROM info_type;
```

We now have a `data.frame` called `info_types` in our environment.

```
# Note: this is an r chunk s we use R instead of SQL.
```

```
head(info_types)
```

	id	info
1	1	runtimes
2	2	color info
3	3	genres
4	4	languages
5	5	certificates
6	6	sound mix

That's all you need to know about how to get data from MySQL into R. The rest of this lab consists of practicing writing SQL queries. It may be useful to reference the full [documentation for SELECT queries](#).

For example, let's say I wanted to find information on the wacky Bill Murray Movie [Groundhog Day](#).

The titles are stored in the `title` field (i.e. column) in the `title` table. [Note: your professor is not responsible for naming these tables and fields!] Each row in the `title` table corresponds to a single movie, but of course, we need to restrict the rows we retrieve to only those where the `title` field equals `Groundhog Day`. The following query achieves this.

Note: SQL does not require the `==` for testing equality, since you aren't ever changing the data.

Note: You have to use `'` single quotes since you are working within a `"` double-quoted string.

In the chunk below we select every column from the title-table where the title-variable equals 'Groundhog Day'

```
SELECT *  
FROM title  
WHERE title = 'Groundhog Day';
```

Table 4: Displaying records 1 - 5

id	title	imdb_	kind_id	production_	imdb_	netic	episode_	season_	episode	series_	year	sum
19605	Groundhog Day	NA	7	2014	NA	G6532	19602	1	2	NA	a5e203197e1aa883f7884eb89e92	
27895	Groundhog Day	NA	7	2008	NA	G6532	27822	1	48	NA	74ebd1bfceb83d4bdae480326d0	
38707	Groundhog Day	NA	7	2011	NA	G6532	387052	1	2	NA	b8d621787a0ea75c76eff6b8a6a8	
38411	Groundhog Day	NA	7	2016	NA	G6532	384073	1	11	NA	bcef9ee95ae2bb82eca4ae234108	
33708	Groundhog Day	NA	7	2016	NA	G6532	337080	1	5	NA	fcaddb8be4ab7b9c5ca9f27bb325	

That retrieved a lot of movies! Let's see if we can refine our query. First, movies (as opposed to TV episodes, etc.) have the `kind_id` value of 1.

```
SELECT *
FROM title
WHERE title = 'Groundhog Day'
AND kind_id = 1;
```

Table 5: 1 records

id	title	imdb_	kind_id	production_	imdb_	netic	episode_	season_	episode	series_	year	sum
36642	Groundhog Day	NA	1	1993	NA	G6532	NA	NA	NA	NA	2f0a563d0b0a1f57a19385de5a87	

Now we have the result that I want.

Imagine that I didn't know the full title of the movie I could soften my query by searching for the phrase `Groundhog` within the title. We can do this using the `LIKE` function along with some wildcards (`%` in SQL).

```
SELECT *
FROM title
WHERE title LIKE '%Groundhog%'
AND kind_id = 1;
```

Table 6: 4 records

id	title	imdb_index	production_year	imdb_rating	metacritic_score	seasons	episodes	series_year	5sum	
366427	Groundhog Day	NA	1	1993	NA	G6532	NA	NA	NA	2f0a563d0b0a1f57a19385de5a87
366427	Groundhog	NA	1	2015	NA	G6532	NA	NA	NA	7b71cb8ae79de1171a71f95d2e50
366427	Groundhog	NA	1	2017	NA	G6532	NA	NA	NA	5e7183dbeb6c28fb6445c4013b2
366427	Groundhog Day	NA	1	2016	NA	G6532	NA	NA	NA	c45dd6456b9787e5f71144d6c3a

Pretend I'm still not sure which of the above four movies is the real Groundhog Day movie I'm interested in, but I'm sure its the first movie that came out. I could put them in order with the code below.

```
SELECT *
FROM title
WHERE title LIKE '%Groundhog%'
      AND kind_id = 1
ORDER BY production_year;
```

Table 7: 4 records

id	title	imdb_index	production_year	imdb_rating	metacritic_score	seasons	episodes	series_year	5sum	
366427	Groundhog Day	NA	1	1993	NA	G6532	NA	NA	NA	2f0a563d0b0a1f57a19385de5a87
366427	Groundhog	NA	1	2015	NA	G6532	NA	NA	NA	7b71cb8ae79de1171a71f95d2e50
366427	Groundhog Day	NA	1	2016	NA	G6532	NA	NA	NA	c45dd6456b9787e5f71144d6c3a
366427	Groundhog	NA	1	2017	NA	G6532	NA	NA	NA	5e7183dbeb6c28fb6445c4013b2

Finally I can select just the three columns I'm interested in. Also notice that I am renaming the title table as t. So I select the columns t.title and t.production_year. This is called creating an alias. The convention is table.variable. This will be useful when joining tables.

```
SELECT t.id, t.title, t.production_year
FROM title AS t
WHERE title LIKE '%Groundhog%'
      AND t.kind_id = 1
ORDER BY t.production_year;
```

Table 8: 4 records

id	title	production_year
3664274	Groundhog Day	1993
3664277	Groundhogs	2015
3664276	Groundhog's Day	2016
3664273	Groundhog	2017

Its the first Groundhog day that came out in 1993 with ID 3664274.

Exercise:

Find the original Ghostbusters in the `title` table.

```
SELECT *
FROM title
WHERE title = 'Ghostbusters'
AND kind_id = 1;
```

Table 9: 3 records

id	title	imdb_index	kind_id	production_year	imdb_id	headlines	episode_runtime	episode_runtime	episode_runtime	episode_runtime	series_year	sum
3644551	ghostbusters	NA	1	2016	NA	G2312	NA	NA	NA	NA	9dc8686712d01fc31d597c8aa189	
3644550	ghostbusters	NA	1	1984	NA	G2312	NA	NA	NA	NA	595c84db94e698e6002097cd7c2f	
3644556	ghostbusters	NA	1	NA	NA	G2312	NA	NA	NA	NA	d5eb40dc8725d1041e0260e7baa	

Now lets consider the name table

```
DESCRIBE name
```

Table 10: Displaying records 1 - 5

Field	Type	Null	Key	Default	Extra
id	int	NO	PRI	NA	auto_increment
name	text	NO	MUL	NA	
imdb_index	varchar(12)	YES		NA	
imdb_id	int	YES	MUL	NA	
gender	varchar(1)	YES		NA	

Exercise:

Find [Andie MacDowell](#)'s id in the `name` table.

Note: that names are listed last name first and separated by a comma (eg Murray, Bill)

```
SELECT * FROM name
WHERE name = 'MacDowell, Andie';
```

Table 11: 1 records

id	name	imdb_id	imdb_id	gender	name_pcode	name_pcode	name_pcode	name_pcode	name_pcode
3479179	MacDowell, NA Andie	NA	f	M2345	A5352	M234	b73def7bf61212e32f9c6bd3c39e1c00		

Joining tables

In the IMDB, the `title` table contains information about movies, the `name` table contains the names of people, the `char_name` table contains information about the names of characters, and the `cast_info` table contains information about which people played which roles in which movies. Linking the tables together is essential in order to extract information from the database.

Since we already know that the ID of *Groundhog Day* is 3664274, we can use that to find all of the cast assignments.

```
SELECT *
FROM cast_info
WHERE movie_id = 3664274;
```

Table 12: Displaying records 1 - 5

id	person_id	movie_id	person_role_id	note	nr_order	role_id
118176	17392	3664274	376	NA	43	1
1764292	232622	3664274	352541	(uncredited)	NA	1
2746369	354870	3664274	191351	NA	13	1
3167452	400710	3664274	140101	NA	34	1
4584838	577860	3664274	98	(uncredited)	NA	1

Note that this returns a list of person-role pairs. `person_id` is unique for the actors and `person_role_id` is unique for the role played.

1. Find all the rows in `cast_info` that correspond to Andie MacDowell as an actress.

We first need to understand the columns that are in `cast_info`.

```
DESCRIBE cast_info
```

Table 13: Displaying records 1 - 5

Field	Type	Null	Key	Default	Extra
id	int	NO	PRI	NA	auto_increment
person_id	int	NO	MUL	NA	
movie_id	int	NO	MUL	NA	
person_role_id	int	YES	MUL	NA	
note	text	YES		NA	

The id from the name table should match the `person_id` from the `cast_info` table. So we can join on that. Notice that the names of the tables are being shorted again as `n` for name and `ci` for `cast_info`.

```
##SAMPLE SOLUTION
```

```
SELECT n.name, ci.role_id
FROM cast_info AS ci
JOIN name AS n ON n.id = ci.person_id
WHERE ci.person_id = 3479179;
```

Table 14: Displaying records 1 - 5

name	role_id
MacDowell, Andie	2
MacDowell, Andie	2
MacDowell, Andie	2
MacDowell, Andie	2
MacDowell, Andie	2

Find all the people who acted in Groundhog Day

Next, we can join the `cast_info` table on the `name` table to recover the names of the actors in Groundhog Day.

```
SELECT n.name, ci.role_id
FROM cast_info AS ci
JOIN name n ON n.id = ci.person_id
WHERE movie_id = 3664274;
```

Table 15: Displaying records 1 - 5

name	role_id
Adler, Roger	1
Blakeslee, Douglas	1
Campbell, Ken Hudson	1
Chaiyabhat, Shaun	1
DeGuide, Tony	1

Note how we have used table aliases to save some typing.

Joining more than two tables

Add the names of the characters she played to the list of Andie MacDowell's roles from the previous exercise.

We have to do a second join to get the names of the roles Andie MacDowell has played.

```
## SAMPLE SOLUTION
SELECT n.name, ci.role_id, cn.name
FROM cast_info ci
JOIN name n ON n.id = ci.person_id
JOIN char_name cn ON cn.id = ci.person_role_id
WHERE ci.person_id = 3479179;
```

Table 16: Displaying records 1 - 5

name	role_id	name
MacDowell, Andie	2	Herself
MacDowell, Andie	2	Herself - Presenter

name	role_id	name
MacDowell, Andie	2	Herself - audience member
MacDowell, Andie	2	Herself
MacDowell, Andie	2	Helen Kalahan

Exercise

Find Annie Davis's full filmography, in chronological order. Include each movie's `title`, `production_year`, and the name of the character that she played.

Hint: You will need the `cast_ing`, `title`, and `char_name` tables.

```
SELECT t.title, t.production_year, cn.name
FROM cast_info ci
JOIN title t ON ci.movie_id = t.id
JOIN char_name cn ON cn.id = ci.person_role_id
WHERE ci.person_id = 3479179
      AND t.kind_id = 1
ORDER BY production_year;
```

Table 17: Displaying records 1 - 5

title	production_year	name
Greystoke: The Legend of Tarzan, Lord of the Apes	1984	Miss Jane Porter
St. Elmo's Fire	1985	Dale Biberman
Sex, Lies, and Videotape	1989	Ann Bishop
		Mullany
Green Card	1990	Brontë
The Object of Beauty	1991	Tina

Exercise

Pull up the titles of the movies that were made in my birth year 1984. There are more than 1000 so we'll put a limit on the number of movies returned.

```
SELECT t.title, t.production_year
FROM title AS t
WHERE production_year = 1984
```

```
LIMIT 100;
```

Table 18: Displaying records 1 - 5

title	production__year
(1984-11-04)	1984
Game 2	1984
(#1.1)	1984
A Winter Harvest	1984
Once a Hero	1984

Exercise

Pull the one movie that Andie MacDowell was in in 1984 and bring it into R as a dataframe.
Hint: We found it above.

```
SELECT t.title, t.production_year
FROM cast_info AS ci
JOIN title AS t ON ci.movie_id = t.id
WHERE t.production_year = 1984
AND ci.person_id = 3479179
AND t.kind_id = 1;
```

Table 19: 1 records

title	production__year
Greystoke: The Legend of Tarzan, Lord of the Apes	1984

Exercise

This [website](#) lists Nicole Kidman as the most prolific actress. Find the name and production year of all of the movies she has been in ordered by production__year descending

```
SELECT * FROM name
WHERE name = 'Kidman, Nicole';
```

Table 20: 1 records

id	name	imdb_index	imdb_gender	name_pcode	name_pcode	name_pcode	name_pcode	name_pcode	name_pcode
3353717	Kidman, Nicole	NA	NA	f	K3524	N2423	K35	69a5531f1e2c7311fcd6c17f1124b397	5

```

SELECT t.title, t.production_year
FROM cast_info AS ci
JOIN title AS t ON ci.movie_id = t.id
WHERE ci.person_id = 3353717
AND t.kind_id = 1
ORDER BY t.production_year DESC;

```

Table 21: Displaying records 1 - 5

title	production_year
Destroyer	2018
Aquaman	2018
Boy Erased	2018
How to Talk to Girls at Parties	2017
The Upside	2017

Exercise

Challenge: Look up the COUNT() and GROUP function for MySQL to answer the question, which year did Nicole Kidman act in the most movies?

```

SELECT t.production_year, COUNT(t.production_year) AS movies
FROM cast_info AS ci
JOIN title AS t ON ci.movie_id = t.id
WHERE ci.person_id = 3353717
AND t.kind_id = 1
GROUP BY t.production_year
ORDER BY movies DESC;

```

Table 22: Displaying records 1 - 5

production_year	movies
2007	8
2015	5
2014	5
2017	5
2006	5